

Programme-1

Load, display, resize, crop, and save images; convert images between color spaces (RGB, grayscale, binary).

1. Load the image

```
img = imread('IMG_20241007_120658.jpg');
```

2. Display

the original

```
figure,  
imshow(img);  
title('Original  
Image');
```

Original Image



3. Resize the image

```
resized_img = imresize(img, [300 400]); % Resize to  
300x400 pixels figure, imshow(resized_img);  
title('Resized Image');
```

Resized Image



4. Crop the image

```
cropped_img = imcrop(img, [50 50 200 200]); % Crop with [x, y,  
width, height] figure, imshow(cropped_img);  
title('Cropped Image');
```

Cropped Image



5. Convert to

Grayscale `gray_img =
rgb2gray(img); figure,
imshow(gray_img);
title('Grayscale Image');`

Grayscale Image



6. Convert to Binary (Black & White)

`bw_img = imbinarize(gray_img); % Convert grayscale to binary
figure, imshow(bw_img);
title('Binary Image');`

Binary Image



Programme- 2

Perform histogram equalization on grayscale images, analyze contrast changes, and plot histograms for comparison.

1. Load the image

```
img = imread('IMG_20241007_120658.jpg');
```

2. Convert to grayscale if not already

```
gray_img = rgb2gray(img);
```

3. Perform histogram equalization

```
equalized_img = histeq(gray_img);
```

4. Compute histograms

```
orig_hist = imhist(gray_img); % Histogram of original image
```

```
eq_hist = imhist(equalized_img); % Histogram of equalized  
image
```

5. Display the original and equalized images

```
figure;  
subplot(2,2,1);  
imshow(gray_img  
g);  
title('Original Grayscale Image');
```

```
figure;  
subplot(2,2,2);  
imshow(equalized_img);  
title('Histogram Equalized Image');
```

Original Grayscale Image



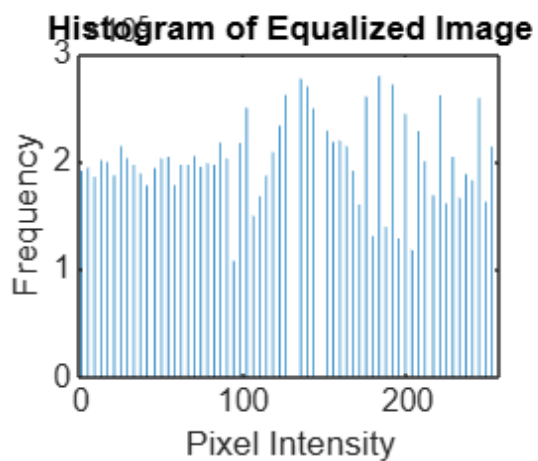

Histogram Equalized Image



6. Plot histograms

```
figure;  
subplot(2,2,3);  
bar(orig_hist);  
title('Histogram of Original Image');  
xlabel('Pixel Intensity');  
ylabel('Frequency');
```

```
figure;  
subplot(2,2,4);  
bar(eq_hist);  
title('Histogram of Equalized  
Image'); xlabel('Pixel Intensity');  
ylabel('Frequency');
```



Programme-3

Apply mean, median, and Gaussian filters to noisy images, comparing results across filter types.

1. Load the image

```
img = imread('IMG_20241007_120658.jpg');  
gray_img = rgb2gray(img); % Convert to grayscale if not already  
figure, imshow(gray_img);  
title('Original Grayscale Image');
```


Original Grayscale Image



2. Add salt & pepper noise

```
noisy_img = imnoise(gray_img, 'salt & pepper', 0.02);  
figure, imshow(noisy_img);  
title('Noisy Image (Salt & Pepper)');
```

Noisy Image (Salt & Pepper)



3. Apply Mean filter (Average filter)

```
mean_filter = fspecial('average', [3 3]);  
mean_filtered = imfilter(noisy_img, mean_filter, 'replicate');  
figure, imshow(mean_filtered);  
title('Mean Filtered Image');
```

Mean Filtered Image



4. Apply Median filter

```
median_filtered = medfilt2(noisy_img, [3 3]);  
figure, imshow(median_filtered);  
title('Median Filtered Image');
```


Median Filtered Image



5. Apply Gaussian filter

```
gaussian_filter = fspecial('gaussian', [3 3], 0.5);  
gaussian_filtered = imfilter(noisy_img, gaussian_filter, 'replicate');  
figure, imshow(gaussian_filtered);  
title('Gaussian Filtered Image');
```

Gaussian Filtered Image



Programme-4

Apply high-pass and Laplacian filters to sharpen image details, analyze before and after images.

1. Load Image and Convert to Grayscale

```
img = imread('IMG_20241007_120658.jpg'); % Replace with your image
figure, imshow(img, []), title('Original Image');
img = rgb2gray(img); % Convert to grayscale if it's RGB
img = double(img); % Convert to double for processing
```


Original Image



2. Fourier Transform and Shift

```
F = fft2(img);
F_shifted = fftshift(F); % Shift zero frequency to center
figure, imshow(log(1 + abs(F_shifted)), []), title('Fourier Transform');
```

Fourier Transform



3. Get Image Size

```
[M, N] = size(img);
u = 0:M-1;
v = 0:N-1;
[U, V] = meshgrid(v - N/2, u - M/2); % Frequency coordinates
```


4. Create High-Pass Filter (HPF) - Circular mask

```
D0 = 50; % Cutoff frequency  
HPF = double(sqrt(U.^2 + V.^2) > D0);  
figure, imshow(HPF, []), title('High-Pass Filter Mask');
```

High-Pass Filter Mask




5. Apply High-Pass Filter in Frequency Domain

```
F_HPF = F_shifted .* HPF;
```

6. Inverse FFT for HPF Image

```
F_HPF_shifted = ifftshift(F_HPF);  
img_HPF = abs(fft2(F_HPF_shifted));  
figure, imshow(img_HPF, []), title('High-Pass Filtered Image');
```

High-Pass Filtered Image



7. Apply Laplacian Filter in Spatial Domain

```
laplacian_filter = [0 -1 0; -1 4 -1; 0 -1 0];  
img_Laplacian = imfilter(img, laplacian_filter, 'replicate');  
figure, imshow(img_Laplacian, []), title('Laplacian Filtered Image');
```

Laplacian Filtered Image



8. Sharpened Image by Adding Laplacian to Original

```
img_Sharpened = img + img_Laplacian;  
figure, imshow(img_Sharpened, []), title('Sharpened Image');
```

Sharpened Image



Programme-5

Use Fourier Transform to switch to frequency domain, apply low-pass and high-pass filters, and then convert back to spatial domain.

1. Read image and convert to grayscale

```
img = imread('IMG_20241007_120658.jpg'); % Replace with your image
figure, imshow(img, []), title('Original Image');
img = rgb2gray(img); % Convert to grayscale if it's RGB
img = double(img); % Convert to double for processing
```

Original Image



2. Fourier Transform and shift


```
F = fft2(img);
F_shifted = fftshift(F); % Centering the frequency components
```

3. Get image size

```
[M, N] = size(img);
u = 0:M-1;
v = 0:N-1;
[U, V] = meshgrid(v - N/2, u - M/2); % Frequency coordinates
```

4. Create Low-Pass Filter (LPF) - Circular mask


```
D0 = 50; % Cutoff frequency  
LPF = double(sqrt(U.^2 + V.^2) <= D0);  
figure, imshow(LP, []), title('Low-Pass Filter Mask');
```



5. Create High-Pass Filter (HPF) - Circular mask

```
HPF = double(sqrt(U.^2 + V.^2) > D0);  
figure, imshow(HPF, []), title('High-Pass Filter Mask');
```

High-Pass Filter Mask



6. Apply filters

```
F_LPF = F_shifted .* LPF; % Low-pass filtered frequency domain  
F_HPF = F_shifted .* HPF; % High-pass filtered frequency domain
```

7. Inverse FFT (Low-Pass)

```
F_LPF_shifted = ifftshift(F_LPF);  
img_LPF = abs(fft2(F_LPF_shifted));  
figure, imshow(img_LPF, []), title('Low-Pass Filtered Image');
```


Low-Pass Filtered Image



8. Inverse FFT (High-Pass)

```
F_HPF_shifted = ifftshift(F_HPF);  
img_HPF = abs(fft2(F_HPF_shifted));  
figure, imshow(img_HPF, []), title('High-Pass Filtered Image');
```

High-Pass Filtered Image



Programme-6

Apply Wiener filter to noisy images, compare restored images with original and noisy versions.

1. Read and convert image to grayscale

```
img = im2double(imread('flower.jpg')); % You can replace with any image
```

2. Add Gaussian noise

```
noisy_img = imnoise(img, 'gaussian', 0, 0.01); % Mean=0, Variance=0.01
```

3. Define filter window size

```
window_size = 5;  
half_window = floor(window_size / 2);
```

4. Pad image to handle border pixels

```
padded_img = padarray(noisy_img, [half_window, half_window], 'symmetric');
```

5. Initialize filtered image

```
filtered_img = zeros(size(noisy_img));
```

6. Compute local mean and variance manually

```
[M, N] = size(noisy_img);  
for i = 1:M  
    for j = 1:N  
        % Extract local region  
        local_window = padded_img(i:i+window_size-1, j:j+window_size-1);  
  
        % Compute local mean and variance  
        local_mean = mean(local_window(:));  
        local_variance = var(local_window(:));  
  
        % Noise variance estimation (Assuming Gaussian noise variance is known)  
        noise_variance = 0.01;  
  
        % Apply Wiener filtering formula  
        filtered_img(i, j) = local_mean + (max(local_variance - noise_variance, 0) /  
local_variance) * (noisy_img(i, j) - local_mean);  
    end  
end
```

7. Display images

```
figure;  
imshow(img), title('Original Image');  
figure;  
imshow(noisy_img), title('Noisy Image');  
figure;
```

```
imshow(filtered_img), title('Manually Filtered Image');
```

Original Image



Noisy Image



Manually Filtered Image



Programme-7

Apply Sobel, Prewitt, and Canny edge detectors, comparing sensitivity and accuracy of edge maps.

1. Read the input image

```
image = imread('image.jpg');  
grayImage = rgb2gray(image);
```

2. Apply edge detection algorithms to generate edge maps

```
cannyEdges = edge(grayImage, 'Canny');  
sobelEdges = edge(grayImage, 'Sobel');  
prewittEdges = edge(grayImage, 'Prewitt');
```

3. Read Ground Truth Image

```
groundTruth = imread('ground_truth.png');
```

4. Perform the comparative analysis for each edge detection method

```
[resultsCanny] = evaluateEdgeDetection(cannyEdges, groundTruth);  
[resultsSobel] = evaluateEdgeDetection(sobelEdges, groundTruth);  
[resultsPrewitt] = evaluateEdgeDetection(prewittEdges, groundTruth);
```

5. Display results

```
fprintf('Canny - Sensitivity: %.4f, Accuracy: %.4f\n', ...  
    resultsCanny.sensitivity, resultsCanny.accuracy);  
fprintf('Sobel - Sensitivity: %.4f, Accuracy: %.4f\n', ...  
    resultsSobel.sensitivity, resultsSobel.accuracy);  
fprintf('Prewitt - Sensitivity: %.4f, Accuracy: %.4f\n', ...  
    resultsPrewitt.sensitivity, resultsPrewitt.accuracy);
```

6. Function to calculate the performance metrics for edge detection

```
function results = evaluateEdgeDetection(detectedEdges, groundTruth)
```

```
% True positives (TP)
```

```
TP = sum(sum(detectedEdges == 1 & groundTruth == 1));
```

```
% False positives (FP)
```

```
FP = sum(sum(detectedEdges == 1 & groundTruth == 0));
```

```
% False negatives (FN)
```

```
FN = sum(sum(detectedEdges == 0 & groundTruth == 1));
```

```
% True negatives (TN)
```

```
TN = sum(sum(detectedEdges == 0 & groundTruth == 0));
```

```
% Sensitivity (Recall or True Positive Rate)
```

```
sensitivity = TP / (TP + FN);
```



```
% Accuracy  
accuracy = (TP + TN) / (TP + FP + FN + TN);
```

```
% Store results in a structure  
results.sensitivity = sensitivity;  
results.accuracy = accuracy;
```

7. Save the generated ground truth images

```
imwrite(cannyEdges, 'ground_truth_canny.png');  
imwrite(sobelEdges, 'ground_truth_sobel.png');  
imwrite(rewittEdges, 'ground_truth_rewitt.png');  
  
end
```

Original Image



Sobel Edges



Prewitt Edges



Canny Edges



```
>> p72
```

```
Canny - Sensitivity: 0.5407, Accuracy: 0.9543
```

```
Sobel - Sensitivity: 0.3314, Accuracy: 0.9847
```

```
Prewitt - Sensitivity: 0.3310, Accuracy: 0.9848
```

(fig. 1) Comparing sensitivity and accuracy of edge maps

Programme-8

Apply global and adaptive thresholding, followed by morphological operations (dilation, erosion) to isolate objects in images.

Read the image

```
img = imread('image.jpg');  
grayImg = rgb2gray(img);
```

Global Thresholding (Otsu's Method)

```
level = graythresh(grayImg); % Compute global threshold  
globalBW = imbinarize(grayImg, level); % Apply threshold
```

Adaptive Thresholding (Local)

```
adaptiveBW = imbinarize(grayImg, 'adaptive', 'Sensitivity', 0.6); % Increased sensitivity
```

Morphological Operations

```
se = strel('disk', 3); % Structuring element (disk shape)
```

Dilation (expands object boundaries)

```
dilatedImg = imdilate(adaptiveBW, se);
```

Erosion (removes small noise)

```
erodedImg = imerode(dilatedImg, se);
```

Additional Morphological Processing

```
closedImg = imclose(erodedImg, se); % Closing to fill small holes
```

Edge Detection

```
edgeImg = edge(grayImg, 'canny'); % Detect edges using Canny method
```

Save output images

```
imwrite(globalBW, 'global_threshold.png');  
imwrite(adaptiveBW, 'adaptive_threshold.png');  
imwrite(dilatedImg, 'dilated.png');  
imwrite(erodedImg, 'final_output.png');  
imwrite(closedImg, 'closed.png');  
imwrite(edgeImg, 'edges.png');
```

Original Image





(fig. 1)global_threshold



(fig. 2)final_output



(fig. 3) dilated



(fig. 4) closed



(fig. 5) adaptive_threshold

Programme-9

Use morphological operations like opening, closing, and boundary extraction to analyze shapes and perform basic object recognition.

1. Load the image

```
img = imread('image.jpg');  
gray_img = rgb2gray(img); % Convert to grayscale if not already  
binary_img = imbinarize(gray_img); % Convert to binary (black & white)
```

2. Define structuring element

```
se = strel('disk', 5); % Circular structuring element with radius 5  
% Apply morphological opening (removes small noise)  
opened_img = imopen(binary_img, se);  
% Apply morphological closing (fills small gaps)  
closed_img = imclose(binary_img, se);  
% Perform boundary extraction (finds object edges)  
boundary_img = binary_img - imerode(binary_img, se);
```

3. Display images separately

```
figure, imshow(binary_img);  
title('Binary Image');  
figure, imshow(opened_img);  
title('After Morphological Opening');  
figure, imshow(closed_img);  
title('After Morphological Closing');  
figure, imshow(boundary_img);  
title('Boundary Extraction');
```

Original Image



Binary Image



After Morphological Opening



Boundary Extraction



Programme-10

Compress an image using DCT, explore effects of compression on image quality, and analyze the trade-offs between quality and file size.

1. Load the original color image

```
img = im2double(imread('image.jpg'));
```

2. Convert image to YCbCr color space (to process luminance separately)

```
ycbcr_img = rgb2ycbcr(img);  
Y = ycbcr_img(:,:,1); % Luminance channel  
Cb = ycbcr_img(:,:,2);  
Cr = ycbcr_img(:,:,3);
```

3. Apply DCT on the luminance channel

```
dct_Y = dct2(Y);
```

4. Define compression levels (percentage of coefficients retained)

```
compression_levels = [0.1, 0.3, 0.5]; % 10%, 30%, and 50% of coefficients retained
```

5. Process images at different compression levels

```
for i = 1:length(compression_levels)  
    thresh = max(dct_Y(:)) * (1 - compression_levels(i)); % Define threshold  
    dct_compressed = dct_Y .* (abs(dct_Y) > thresh); % Zero out small coefficients  
    Y_reconstructed = idct2(dct_compressed); % Apply inverse DCT to reconstruct the  
    luminance channel  
  
    % Reconstruct the color image in YCbCr space  
    compressed_img = ycbcr_img;  
    compressed_img(:,:,1) = Y_reconstructed; % Replace the Y channel  
  
    % Convert back to RGB color space  
    final_img = ycbcr2rgb(compressed_img);  
  
    % Display and save results  
    figure, imshow(final_img);  
    title(['Reconstructed Image (' , num2str(compression_levels(i) * 100), '% coefficients  
    retained')]);
```

6. Save compressed image


```
filename = ['compressed_' num2str(compression_levels(i) * 100) '.jpg'];  
imwrite(final_img, filename);
```

```
end
```

Original Image



Reconstructed Image (10% coefficients retained)



Reconstructed Image (30% coefficients retained)



Reconstructed Image (50% coefficients retained)

